

Disentangling Predictive Surprisal, Entropy, and Syntactic Integration Cost in Human Sentence Processing

Mid-Project Evaluation

Harshit Gupta (2022101124) Sanchit Jalan (2022101070)

24 April 2026

Outline

- 1 Motivation
- 2 Data & Pipeline
- 3 Results
- 4 Summary

The Core Question

Why do some words take longer to read than others?

Predictive Processing

Unexpected words are harder to process.

$$S(w_i) = -\log_2 P(w_i \mid \text{context})$$

Hale 2001; Levy 2008

Dependency Locality Theory

Structurally distant words are harder to integrate.

$$IC(w_i) = |\text{pos}(w_i) - \text{pos}(\text{head}(w_i))|$$

Gibson 2000; Demberg & Keller 2008

Our Contribution

Test **both simultaneously** using Transformer LMs + Bayesian hierarchical modeling across 3 architectures and 181 readers.

Dataset: Natural Stories Corpus

- ~10,000 words of English narrative text
- **181 participants**, self-paced reading times
- **835,700 total observations** after cleaning


Pre-processing:

- RT filter: 100 ms – 3000 ms
- Per-subject $\pm 3\sigma$ outlier removal
- Dependent variable: $\log(\text{RT})$
- Universal Dependencies parses via Stanza

Scale

Item	Count
Words (unique positions)	10,245
Subjects	181
Sentences	485
Total observations	835,700
After cleaning	835,700

Full Pipeline — All 6 Steps Complete



✓ All 6 pipeline steps complete. Results are in hand.

Bayesian Hierarchical Model Specification

$$\log(\text{RT}_{ij}) = \alpha + u_{0i} + \sum_k (\beta_k + u_{ki}) X_{kj} + \varepsilon_{ij}$$

Priors:

$$\begin{aligned}\alpha, \beta_k &\sim \mathcal{N}(0, 1) \\ \sigma_{u_0}, \sigma_{\text{slope}} &\sim \text{HalfNormal}(0.5) \\ \sigma_{\varepsilon} &\sim \text{HalfNormal}(0.5)\end{aligned}$$

9 model variants fitted:

- baseline (trigram only)
- deep_gpt2 / bert / t5
- surprisal_vs_ic
- surprisal_vs_entropy $\times 3$
- full (all predictors)

Key feature


Random slopes for GPT-2 surprisal and IC **per subject** — enables quantifying individual reader strategies (H6).

H1: Deep Neural Surprisal $\gg$ Shallow N-gram


Model	ELPD	β	Sig.
GPT-2	-1163	+0.010	✓
T5	-1191	+0.013	✓
BERT	-1210	+0.005	✓
Trigram	-1211	-0.003	✓

Finding

GPT-2 outperforms the trigram baseline by **48 ELPD units**. Human reading relies on long-distance context, not just local collocations.



H4: Autoregressive Architecture Best Mirrors Human Reading




ELPD Ranking

- 1 **GPT-2** (autoregressive) –1163
- 2 T5 (encoder-decoder) –1191
- 3 BERT (bidirectional) –1210
- 4 Trigram baseline –1211

Why BERT fails

Bidirectional context is unavailable during sequential human reading. A stronger LM is not necessarily a better psycholinguistic model.

H2: Integration Cost Explained Away by Neural Surprisal




Predictor	β	95% HDI	Sign
GPT-2 surprisal	+0.014	[+0.011, +0.018]	✓
Integration cost	-0.001	[-0.003, +0.002]	✗

Finding

The posterior for β_{IC} spans zero symmetrically. Neural surprisal **fully subsumes** the variance previously attributed to structural memory load.

H3: Surprisal and Entropy are Partly Independent

Hyp 3: Surprisal vs. Entropy — independent contributions per architecture



Arch.	β_{surp}	β_{entr}
GPT-2	+0.008 ✓	+0.005 ✓
T5	+0.013 ✓	+0.003 ✓
BERT	+0.005 ✓	+0.000 ✗


Interpretation

GPT-2 & T5: knowing the model was *uncertain* predicts reading difficulty *beyond* just the specific word's probability.

BERT entropy is non-informative — bidirectional pseudo-surprisal conflates the two measures.

H5: Specific Attention Heads Track Syntactic Structure



Architecture	Best Head	ρ	p -value
BERT	L0 H10	−0.624	$< 10^{-100}$
T5	L0 H10	−0.579	$< 10^{-100}$
GPT-2	L5 H9	+0.351	$< 10^{-100}$



Convergent finding

The same head (L0 H10) dominates in both BERT and T5, despite different training objectives — a convergent structural inductive bias.


H5 Deep Dive: BERT L0 H10 Head-Finding Accuracy



Interpretation

Head-finding accuracy: $\approx 60\%$ for $DL=1$, drops to $\approx 15\%$ for $DL>5$ (near chance). Layer 0 has the *highest* mean attention to syntactic heads, declining in deeper layers — a local structural bias that mirrors the human working memory constraint.

H6: Individual Reader Strategies



Component	σ	95% HDI	Sig.
Surprisal slope	0.013	[0.009, 0.016]	✓
IC slope	0.003	[0.000, 0.007]	✓

Finding

$\sigma_{\text{slope}} > 0$: genuine between-reader variation in surprisal sensitivity. Some readers slow sharply at unexpected words; others are barely affected. This heterogeneity is invisible to standard LMER models.

Summary: All 6 Hypotheses Tested

Hyp.	Status	Key Finding
H1: Deep > Shallow	Supported	GPT-2 Δ ELPD = +48 over trigram
H2: Surprisal subsumes IC	Supported	β_{IC} HDI spans zero
H3: Surprisal $\neq$ Entropy	Partial	GPT-2 & T5 yes; BERT entropy n.s.
H4: Autoregressive best	Supported	GPT-2 > T5 > BERT $\approx$ trigram
H5: Heads track syntax	Strongly	$\rho = -0.624$ (BERT L0 H10)
H6: Individual differences	Supported	$\sigma_{\text{slope}} = 0.013$, credibly > 0

5 of 6 fully supported · 1 partially supported

Completed

- Full 6-step pipeline
- GPT-2, BERT, T5 surprisal & entropy
- T5 attention head analysis (new)
- 9 Bayesian model variants
- LOO-CV model comparison
- All 6 hypotheses tested & visualised
- Individual reader slope analysis

Remaining for Final

- 1 **GECO corpus** — cross-dataset validation on eye-tracking data (54K words)
- 2 **MCMC robustness** — address intercept mixing for hierarchical models

Key Takeaways

- 1 **Neural LMs $\gg$ N-grams** for predicting reading difficulty — long-distance context is essential.
- 2 **Architecture modality matters** — left-to-right GPT-2 best mirrors the incremental human reading process.
- 3 **Structural memory load is not independently necessary** — neural surprisal already encodes dependency-based difficulty.
- 4 **Transformers implicitly learn syntax** — the same head (L0 H10) tracks dependency structure in both BERT and T5, suggesting a convergent structural inductive bias.
- 5 **Reading is not one-size-fits-all** — Bayesian modeling reveals meaningful individual differences in predictive processing that standard LMERS miss.

- ① Hale, J. (2001). A probabilistic Earley parser as a psycholinguistic model. *NAACL*.
- ② Levy, R. (2008). Expectation-based syntactic comprehension. *Cognition*, 106(3).
- ③ Demberg, V. & Keller, F. (2008). Data from eye-tracking corpora as evidence for theories of syntactic processing complexity. *Cognition*, 109(2).
- ④ Oh, B. & Schuler, W. (2023). Transformer-based Language Models Explain Human Reading Times. *CogSci*.
- ⑤ Wilcox, E. G. et al. (2020). On the Predictive Power of Neural Language Models for Human Real-Time Comprehension Behavior. *CogSci*.
- ⑥ Clark, K. et al. (2019). What Does BERT Look At? An Analysis of BERT's Attention. *BlackboxNLP*.